




Audian

User & Technical Manual

Voice dictation for Windows — speak naturally and get polished text wherever you type. Private by default: everything runs on your PC.



Contents


- | | | | |
|----------|--------------------------------------|-----------|-------------------------------|
| 1 | What Audian is | 8 | Resource usage (measured) |
| 2 | Installing & first run | 9 | Models: what we chose and why |
| 3 | Dictating | 10 | How it works inside |
| 4 | Knowing when you're done (auto-stop) | 11 | Privacy & your data |
| 5 | Rewriting modes & providers | 12 | Troubleshooting |
| 6 | Per-app modes & translation | 13 | Configuration reference |
| 7 | The Audian window | 14 | Building & publishing |

1 What Audian is

Audian is a small background utility. You press a keyboard shortcut in any app, speak the way you talk, and Audian inserts a clean, well-formed version of what you *meant* — not a literal transcript — at your text cursor.

Saying *"um so basically I, I want to make this thing work better and like organize the code because it's kinda messy"* produces **"I want to make this work better and organize the code, because it's kind of messy."** in Natural mode, or **"Improve the implementation and organize the code to make it cleaner and more maintainable."** in AI Prompt mode.

The pipeline



Typically 0.3 – 1.5 s from the end of your speech to text on screen (CPU only)

Design principles

- **Private by default.** Speech recognition and rewriting run on your CPU. Nothing is sent anywhere unless you choose the Antigravity cloud provider.
- **Invisible until needed.** The tray app idles at ~11 MB of RAM and 0 % CPU. The AI engines are separate processes that start when you press the shortcut and exit as soon as your text is inserted.
- **Works everywhere you type.** Text is pasted through the clipboard, which every Windows app understands — editors, browsers, chat apps, Office, terminals — and your clipboard is restored afterwards.
- **Yours to style.** Six colour themes — Ivory (off-white, the default), Violet, Ocean, Mint, Sunset and Rose — for the window, the logo and the recording indicator, which can also be dark or light.
- **Fails gracefully.** If the AI fails, a basic cleanup is inserted; if the window closed or is elevated, the text goes to your clipboard and a notification explains why.

11 MB

tray app RAM while idle
(0 % CPU)

0.3–1.5 s


end of speech → text,
CPU only

0 MB

GPU memory used —
runs on any PC

2 Installing & first run

Run `Audian-Setup-1.0.0-x64.exe` . It installs for your Windows account only, so no administrator rights are needed (the same model as VS Code's user installer). Requirements: Windows 10 or 11, 64-bit x86 CPU with AVX2 (any Intel/AMD CPU from roughly 2015 on), 8 GB RAM recommended (4 GB works with the lighter models).



The installer. Options: location, start with Windows, desktop shortcut, and downloading the AI models (about 1.5 GB) during setup.

What the installer does


Step	Details
Files	<code>%LOCALAPPDATA%\Programs\Audian</code> : <code>audian.exe</code> (tray app + window), <code>audian-stt.exe</code> (speech engine), <code>audian-llm.exe</code> (rewriting engine), <code>uninstall.exe</code> . About 29 MB.
Shortcuts	Start Menu entry "Audian"; optional desktop shortcut.
Registration	Listed in <i>Settings</i> > <i>Apps</i> > <i>Installed apps</i> with a working Uninstall button.
Startup	Optional entry in your user's <i>Run</i> key (Audian starts in the tray when you sign in).
Models	Optional download of Whisper Small (264 MB) and Qwen 3.5 2B (1.28 GB) into <code>%LOCALAPPDATA%\Audian\models</code> , verified by SHA-256. Resumable.
Updates	Running a newer installer stops Audian, replaces the files and keeps your settings, history and models.

Silent install (for scripts, winget, etc.)

```
Audian-Setup.exe --quiet [--with-models] [--no-autostart] [--desktop-shortcut] [--no-launch]
uninstall.exe --uninstall [--quiet] [--purge]      # --purge also deletes settings, history, models
```

First run

On first launch Audian opens a short welcome: (1) download the models if they're missing, (2) pick and test your microphone, (3) learn the shortcut. After that it lives in the system tray.



The first-run welcome.

3 Dictating

Click into any text field, then use your shortcut (default **Ctrl** + **Shift** + **Space**; yours is set to **Ctrl** + **Shift** + **V**).

Style	How it works	Best for
Hybrid (default)	Hold the shortcut while speaking and release to finish — or tap it once and just talk; Audian finishes when you pause (or when you tap again).	Everything
Push-to-talk	Records only while the shortcut is held.	Short phrases, noisy rooms
Toggle	Press to start; press again or pause to finish.	Long, hands-free dictation


Esc cancels a dictation at any point. **Ctrl** + **Shift** + **Alt** + **V** pastes your last dictation again (configurable).

The recording indicator


A small pill appears near your text cursor (or at the bottom of the screen if the app doesn't expose its cursor). It never takes focus and clicks pass through it.



Listening — live waveform; label shows the mode. A ring around the dot means hands-free.



Pause detected — a line in your theme's colour fills as auto-stop approaches; speaking again resets it.



Done — word count, then the pill fades. Errors show an amber mark and a short reason.

While processing, the pill shows *Transcribing...* then *Polishing...*. Soft chimes mark start and stop (can be turned off under *Microphone & sound*). Its colours, dark or light style, position and label are set under *Appearance*.

Where the text goes

- Audian remembers the window you were in when you pressed the shortcut, and pastes there — even if you clicked elsewhere while it was processing (it switches back first).
- Your clipboard is saved before pasting and restored ~0.6 s later. Dictated text is excluded from **Win** + **V** clipboard history and cloud clipboard.
- In terminals, multi-line results are joined into one line so nothing runs by accident. mintty/PuTTY receive **Shift** + **Insert**.
- If the window was closed, or runs as administrator (Windows blocks input from normal apps into elevated ones), the text is placed on the clipboard and a notification tells you to press **Ctrl** + **V**.

4 Knowing when you're done

In hands-free dictation Audian decides when you've finished speaking — quickly when you've completed a sentence, patiently when you pause mid-thought.

How it decides

1. **Voice activity detection** runs on every 10 ms of audio and tracks your room's background noise, so it knows the difference between “quiet” and “speaking” on any microphone.
2. When you pause (≈ 0.5 s), Audian quietly **transcribes what you've said so far** (a “preview”).
3. If the preview sounds **finished** (“...send it to Maria.”), Audian stops after your *pause tolerance* (default 2.0 s).
4. If it sounds **unfinished** — it ends with “and”, “but”, “because”, “the”, a comma, “um”, or Portuguese “e”, “mas”, “porque”, “que”, “de”... — Audian waits the *extra patience* as well (default +1.5 s).
5. When it stops, the preview is **reused**, so your text usually appears 0.3–1 s later.

Tuning it

Shortcut & recording › Finish automatically. Presets: **Quick** 1.2 s, **Balanced** 2 s, **Patient** 3 s, or any value 0.6–6 s. You can also disable smart detection (pure silence timer) or auto-stop entirely (then tap the shortcut to finish). If nothing is said for 8 s, the dictation is cancelled.

Measured behaviour (test recordings)


What was said	Result
“Let's meet on Wednesday at three.” → silence	Stopped 1.5–2.0 s after the sentence; text appeared 0.98 s later (transcript reused).
“I think we should push the release and” → 2.2 s pause → “also update the documentation before Friday.”	Did not stop at the pause (sentence unfinished); captured everything; stopped 2.07 s after the end; text 0.69 s later.
“Send the report to Maria.” → 2.2 s pause → “Also call John.”	Stopped after the first sentence (complete + pause longer than the tolerance). Choose <i>Patient</i> if you take long breaks between sentences.

5 Rewriting modes & providers

Modes

Mode	Purpose	Example output (from real tests)
Natural	Clean speech, stay close to your words.	"Let's meet on Wednesday at 3. Can you send the report to Maria before that?" (from "Let's meet on Tuesday — no, wait, Wednesday...")
AI Prompt	A clear, structured instruction for an AI assistant.	"Add a settings page where users can select their microphone and choose a shortcut. Include a toggle to enable or disable Windows startup, and ensure all settings are saved to a config file and loaded on startup."
Professional	Polished wording for work email/chat.	"For the new feature, please add a settings page where users can select their microphone, configure a shortcut, and toggle launching on Windows startup..."
Document	Prose for documents, with paragraphs.	—
Literal	Almost verbatim; no AI.	Whisper's transcript with recognizer artefacts removed.
Custom	Only your custom instructions.	e.g. "Always organize my spoken instructions into concise technical prompts."

In every mode the rewriter **never answers** what you say: "uh what's the, what's the capital of France" becomes "What's the capital of France?", not "Paris". It keeps your language (Portuguese stays Portuguese), your point of view, names, numbers and code identifiers.



Rewriting: default mode, AI provider and options.

Providers


Provider	Where text goes	Speed	Notes
Local AI default	Nowhere — runs on your CPU	0.4–1 s	Qwen 3.5 2B; private and offline.
Antigravity cloud	Your Antigravity account (Gemini/Claude models)	4–8 s	Best on long, rambling dictation. Audio never leaves the PC — only text. Each dictation appears as a conversation in your Antigravity history.
Offline rules	Nowhere	instant	Removes “um/uh”, stutters (“I I”), fixes spacing/capitals. Also the automatic fallback when an AI provider fails.

6 Per-app modes & translation

Per-app modes

Audian checks which app you're dictating into and picks the mode for it automatically. Rules match the program's executable name or a word in its window title (so "ChatGPT" in a browser tab works too).

Profile (default)	Mode	Matches
AI assistants	AI Prompt	claude.exe, ChatGPT.exe; titles containing ChatGPT, Claude, Gemini, Copilot, Perplexity
Code editors	AI Prompt	Code.exe, Cursor.exe, Windsurf.exe, Antigravity IDE.exe, zed.exe, devenv.exe, idea64.exe, pycharm64.exe
Email	Professional	OUTLOOK.EXE, olk.exe, thunderbird.exe; titles Gmail, Outlook
Documents	Document	WINWORD.EXE, soffice.bin, Notion.exe, Obsidian.exe; title Google Docs
Work chat (off)	Professional	ms-teams.exe, Teams.exe, slack.exe



Per-app modes are fully editable; the overlay label shows which mode is active.

Output language (translate as you dictate)


Under *Rewriting* › *Output language* (or the tray menu) choose a language to always write in. Speak Portuguese, get English — or the reverse. Tested results: "Você pode me explicar como funciona o borrow checker do Rust?" → "Can you explain how the Rust borrow checker works?"; "Can you check if the build is passing before we merge?" → "Pode verificar se o build está passando antes de fazer o merge?". English and Portuguese have built-in examples for the local model; other target languages work best with the Antigravity provider.

Recognizing two languages

With *Spoken language: Auto-detect*, tick the languages you actually speak (e.g. English + Portuguese). Detection then only chooses among those — avoiding mix-ups such as short Portuguese phrases being detected as Galician or Spanish — and is faster.


7 The Audian window

Left-click the tray icon to open it (right-click shows a quick menu with mode, provider, output language, microphone, auto-stop and per-app toggles). All changes are saved and applied instantly.




Home: status, your shortcut, statistics, quick settings, live resource usage and recent dictations.


Page	What you'll find
Home	Status, words dictated, time saved vs typing, quick toggles, live engine usage, recent dictations.
Shortcut & recording	Shortcut editor with conflict check, recording style, auto-stop tuning, Esc-to-cancel, paste-last shortcut.
Microphone & sound	Device, live level meter, gain, voice-detection sensitivity, silence threshold, max length, sound cues.
Speech recognition	Whisper model, CPU threads, how long the model stays loaded, languages, custom vocabulary.
Rewriting	AI on/off, default mode, provider, output language, custom instructions, fallback.
Per-app modes	Automatic mode rules.
Models	Download, switch and delete models; sizes, RAM and speed for each.
History	Searchable recent dictations with what you originally said, timing and provider; copy or clear.
General	Start with Windows, notifications, insertion method, clipboard restore, terminals.
Appearance	Colour theme (window, logo, indicator); recording indicator style, position and label, with a live preview.
Privacy	Data-flow summary, history, recordings, transcript logging, statistics.
About	Version, live resource usage, file locations.




Shortcut & recording: the shortcut editor checks availability as you type.



Appearance: six themes that apply instantly, and the recording indicator in dark or light style.



Speech recognition: model and how long it stays in memory — by default it exits right after each dictation.



Models: download, switch and remove, with size, RAM and speed for each.

8 Resource usage (measured)

All numbers were measured on the development PC: Intel Core i9-12900KS (16 cores / 24 threads), 16 GB RAM, Windows 11. The GPU (RTX 5060) is **not used**: VRAM usage is zero.

The tray app

State	RAM (working set)	CPU
Idle (waiting for the shortcut)	11 MB (1.9 MB private)	0 ms in 30 s — no timers or polling run while idle
Recording (15.7 s dictation)	~38 MB (audio buffer, overlay)	0.42 s total = 0.11 % of the machine
Right after a dictation	1.5 MB (memory handed back to Windows)	0
Audian window open	~150 MB (separate process, OpenGL)	only while animating

AI engines (separate processes)

Engine / model	Load time	Peak RAM	CPU while working	Speed	VRAM
Speech: Whisper Small Q8 default	0.45 s + 0.25 s warm-up (while you speak)	~450 MB	8 threads (~7 cores, ~30 % of this CPU)	6.6 s clip: 0.4–0.6 s 14.6 s clip: 1.0–1.4 s	0
Speech: Whisper Large v3 Turbo Q5	2.2 s	777 MB	~7.6 cores	14.6 s clip: 4.4–7.9 s	0
Rewriting: Qwen 3.5 2B Q4 default	0.5 s + 25 ms prompt restore from cache (while you speak)	1.4 GB 0.5 GB private	~5.4 cores	0.35–1.0 s per rewrite (~30 tokens/s)	0
Rewriting: Antigravity	CLI start ≈4 s (hidden while you speak)	not measured (CLI process)	negligible (runs in the cloud)	4–8 s	0


By default both engines **exit as soon as your text is inserted** ("Keep the model in memory: After use"), so between dictations Audian uses only the tray app's ~11 MB. Because the engines start the moment you press the shortcut and load while you speak, this costs almost nothing: in a measured dictation starting from nothing loaded, the text was inserted 1.9 s after the recording ended (0.6 s speech recognition + 1.2 s rewriting). If you dictate in quick bursts, 1 min – 1 h keeps them loaded in between, at 0 % CPU.

Where the memory goes

While a dictation is processed	Memory	Notes
Rewriting model weights (Qwen 3.5 2B)	~0.9 GB	Memory-mapped straight from the model file: shared, file-backed pages that Windows can drop at any time. They stay in the file cache after the engine exits, which is why reloading takes only 0.5 s.
Rewriting engine, private	~0.5 GB	Compute buffers, attention and recurrent state, the ~27 MB prompt snapshot, the 248k-token vocabulary.
Speech engine (Whisper Small Q8)	~0.4 GB	Model weights (264 MB) plus encoder/decoder buffers.

Task Manager's default *Memory* column counts only private memory, so it shows less than Audian's own resource card, which shows each engine's full working set. Version 0.2.0 peaked at 1.96 GB for the rewriting engine: llama.cpp kept a second, "repacked" copy of most weights (0.5 GB) and reserved a 0.5 GB compute buffer for 512-token batches. Audian now uses the mapped weights directly and 128-token batches, with no measurable change in rewriting speed.

Only working when necessary



- The microphone is opened only while recording (the Windows mic indicator is off otherwise).
- The engines are separate processes started by the shortcut, not at boot, and they exit as soon as the dictation is done — also when it is cancelled or contained no speech — or after the idle time you chose. A Windows Job Object guarantees they never outlive Audian.
- The tray app uses no timers while idle: key-release polling runs only while you hold the shortcut, voice detection only during hands-free dictation, and the overlay animation only while it is visible.
- Preview transcriptions happen only during pauses in hands-free mode.
- OpenMP is disabled in the rewriting engine, so its threads don't spin-wait after each request.
- The settings window is a separate process, so its GUI memory exists only while it is open.
- After each dictation the tray app returns its temporary memory to Windows.

9 Models: what we chose and why

Speech recognition — Whisper Small (Q8) by default

Whisper (OpenAI, MIT licence) running through `whisper.cpp` is the most mature offline speech engine on Windows, handles 99 languages including Portuguese, and needs no GPU. On a CPU the decisive factor is speed:

Model	Disk	RAM	Speed on this PC	Portuguese error rate*	Choose it when...
Whisper Base	148 MB	~250 MB	≈0.2 s	~13.9 %	old or low-power PCs
Whisper Small Q8	264 MB	456 MB	0.4–1.4 s	~7.7 %	default — best speed/accuracy on CPU
Whisper Small English	264 MB	456 MB	≈0.4 s	—	you only speak English
Whisper Medium Q5	539 MB	~800 MB	≈2–4 s	~5.1 %	middle ground
Whisper Large v3 Turbo Q5	574 MB	777 MB	3–8 s	~4.2 %	accents, noisy audio; or with a GPU build

* Word error rate on the FLEURS Portuguese test set (public benchmark at Q8, lower is better).

Why Small: it is 4–6× faster than Large v3 Turbo on this CPU while being good enough for clean dictation — and the rewriting step fixes most remaining recognition slips (“end” → “and”). Turbo is the better choice once speed is not limiting (e.g. a GPU build). **Worth adding next:** NVIDIA *Parakeet TDT 0.6B v3* (ONNX) is roughly 2× faster than Small on CPU with Turbo-like accuracy (~5 % Portuguese error rate) for 25 European languages — it would plug in behind the same `TranscriptionProvider` interface but requires the ONNX Runtime.

Rewriting — Qwen 3.5 2B (Q4_K_M) by default

Model	Disk	RAM	Rewrite time*	Strengths / limits
Qwen 3.5 0.8B Q4	533 MB	~0.7 GB	≈0.3 s	Lightest; fine for cleanup, weaker at restructuring (AI Prompt).
Qwen 3.5 2B Q4	1.28 GB	1.4 GB	0.4–1 s	Default. Strong instruction following for its size, 100+ languages with good Portuguese, Apache-2.0.
Qwen 3.5 4B Q4	2.74 GB	~2.9 GB	≈1.5–2.5 s	Better on long, rambling dictation; needs 16 GB RAM.
Llama 3.2 3B Instruct Q4	2.02 GB	~2.5 GB	≈1–2 s	Good English writing; weaker in Portuguese.
Gemma 4 E2B Q4 (not in catalog)	3.1 GB	~3.4 GB	≈1.5 s	Strong multilingual; large file for its speed.
SmolLM3 3B / Phi-4-mini 3.8B (not in catalog)	1.9 / 2.5 GB	2.3 / 2.9 GB	≈1–2 s	Good English; fewer languages.
Antigravity (cloud)	0	not measured	4–8 s	Best quality; needs internet; text leaves the PC.

* Typical sentence on this CPU. RAM is the engine's working set while rewriting, including the memory-mapped model file; figures with “~” are estimates from file size and context shape.

Why Qwen 3.5 2B: rewriting needs a model that follows rules precisely (keep the meaning, don't answer questions, keep the language) while being fast enough to feel instant. In our tests the 2B model resolved self-corrections, removed fillers, kept Portuguese as Portuguese, never answered questions and translated reliably, at under a second per rewrite on the CPU — with about 1.3 GB of disk. The 0.8B model is noticeably weaker at AI-Prompt restructuring; 4B models are better on long inputs but 2–3× slower and need twice the RAM.

Using your GPU (future option)

Both engines can be built with GPU backends (Vulkan or CUDA for the RTX 5060). That would make Whisper Large v3 Turbo take ~0.3 s and allow bigger rewriting models, at the cost of ~1–3 GB of VRAM while loaded. Audian ships CPU-only so it works on every PC with no drivers involved.

Any whisper.cpp `ggml-*.bin` model or instruction-tuned GGUF chat model placed in the models folder appears in the app as a “custom” model.

10 How it works inside

Process model

Process	Role	Technology
<code>audian.exe</code> (tray)	Hotkeys, recording, voice detection, overlay, tray, insertion, orchestration	Rust; raw Win32 message loop; WASAPI via cpal; tiny-skia + fontdue rendering
<code>audian.exe --settings</code>	The Audian window	egui (OpenGL), separate short-lived process
<code>audian-stt.exe</code>	Speech recognition	whisper.cpp (ggml, AVX2)
<code>audian-llm.exe</code>	Local rewriting	llama.cpp (ggml, AVX2, no OpenMP)
<code>agy.exe</code> (optional)	Cloud rewriting through Antigravity	Antigravity CLI with a tool-less custom agent

Engines talk to the tray app over their stdin/stdout pipes with a framed protocol (length-prefixed JSON control messages, binary frames for 16 kHz PCM audio). Separate processes keep the tray app tiny, return model memory to Windows on exit, isolate crashes in native code, and let whisper.cpp and llama.cpp each use their own ggml library (they cannot be linked into one binary).

A dictation, step by step

1. **Shortcut** (`RegisterHotKey` ; conflicts are reported). Audian notes the focused window, its caret position, process and title (title is used only for per-app modes, never logged).
2. **Recording** starts within ~20 ms (WASAPI shared mode, device's native rate). In parallel, the engines start and load their models; the rewriting engine restores its prompt prefix from the cache.
3. **Voice detection** on 10 ms blocks; the overlay animates from the same levels.
4. **End**: key release (push-to-talk), second tap, or auto-stop.
5. **Audio preparation**: windowed-sinc resampling to 16 kHz, trimming of leading/trailing silence, normalisation.
6. **Transcription** (unless the pause preview can be reused).
7. **Rewriting** with the session's mode; output sanitised (strip echoed tags/labels/quotes, reject outputs far longer than the input — the sign of an answer instead of a rewrite).
8. **Insertion**: verify/restore focus → wait for modifier keys to be released → clipboard snapshot → paste → restore clipboard after 0.6 s unless you copied something new meanwhile.
9. **History & stats** recorded (text only); engines exit (with “After use”); memory trimmed.

Voice activity detection

Each 10 ms block's loudness (dBFS) is compared with an adaptive noise floor. The floor follows the quietest recent level — falling quickly, rising only 3 dB/s — so steady noise like fans is learned within seconds while the natural dips between words keep it from creeping up during speech. A block is voiced when it is *margin* dB above the floor (16 dB at sensitivity 0, 6 dB at sensitivity 1; default 11 dB; never below -60 dBFS). Speech starts after 3 voiced blocks (30 ms: ignores clicks) and ends after 200 ms of unvoiced blocks (bridges gaps between words). The detector also rejects recordings with no speech before any AI runs.

Speech recognition optimisations

Technique	Effect (measured)
Compiler optimisation + AVX2 forced for whisper.cpp under MSVC	21.4 s → 2.4 s for a 6.6 s clip (the default MSVC build had no optimisation)
Encoder context sized to the audio length instead of a fixed 30 s window	2.4 s → 0.37 s
Warm-up inference while you speak	removes a ~1.5 s first-dictation penalty
Language detection restricted to your languages	avoids misdetection on short phrases
Glossary prompt from your custom vocabulary	favours your spellings of names and jargon

Local rewriting

- **Prompt:** strict rules (edit, never answer; keep language, meaning, names; remove fillers; resolve self-corrections) + a mode style + your custom instructions + 6 few-shot examples written like real Whisper output, in English and Portuguese.
- **Prefix caching:** the system prompt and examples (~640 tokens) are evaluated once and the model state is snapshotted (~27 MB), in memory and on disk (`%LOCALAPPDATA%\Audian\cache` , the 6 most recent prompts). A freshly started engine restores it in ~25 ms instead of re-evaluating it (~3.5 s), and each dictation only processes its own ~40–60 tokens. Qwen 3.5 is a hybrid recurrent model, so a state snapshot is used rather than truncating a KV cache.
- **Lean memory:** weights are used straight from the memory-mapped file (llama.cpp's repacked weight copy is disabled) and the micro-batch is 128 tokens, which keeps the compute buffer — sized by the vocabulary's logits — at ~0.13 GB instead of ~0.5 GB.
- **Decoding:** greedy (deterministic), “thinking” disabled, output length capped relative to the input.
- **Translation** uses example answers in the target language, which a 2B model follows far more reliably than an instruction alone.

Antigravity integration


Audian keeps a private workspace (`%LOCALAPPDATA%\Audian\agy-workspace`) containing a custom agent with no tools and a minimal prompt — about 700 input tokens per request instead of ~12 000 for the default coding agent. The CLI is started in streaming mode when you press the shortcut, so its start-up overlaps your speech.

Text insertion details

- Clipboard snapshot copies every memory-based format (text, rich text, HTML, images as DIB, file lists), up to 64 MB.
- Dictated text is tagged `ExcludeClipboardContentFromMonitorProcessing` , `CanIncludeInClipboardHistory=0` , `CanUploadToCloudClipboard=0` .
- For Alt/Win shortcuts, an unassigned key is injected so releasing the modifier doesn't open menus or Start.
- Elevation is detected by comparing token integrity levels; elevated targets receive the text via clipboard + notification.
- “Simulate typing” sends Unicode key events (newlines as Shift+Enter so chat apps don't send early); used as fallback when the clipboard is locked.

11 Privacy & your data

Data	Where it goes	Stored?
Microphone audio	Memory of the tray app → speech engine (local pipe)	No. Optional “Keep audio recordings” saves WAV files (off by default).
Transcript & final text	Local engines; or Antigravity if you selected it	History (text only, last 200, can be turned off/cleared). Not logged unless you enable transcript logging.
Window title	Used in memory to match per-app modes	Never stored or logged.
Statistics	Counters only (dictations, words, seconds)	<code>stats.json</code> ; resettable.



The Privacy page shows the data flow for your current settings.

Files

What	Location
Program	<code>%LOCALAPPDATA%\Programs\Audian</code>
Settings	<code>%APPDATA%\Audian\config.toml</code>
Models	<code>%LOCALAPPDATA%\Audian\models</code>
History, statistics	<code>%LOCALAPPDATA%\Audian\history.jsonl</code> , <code>stats.json</code>
Logs	<code>%LOCALAPPDATA%\Audian\logs</code> (rotated at 2 MB)
Prompt cache	<code>%LOCALAPPDATA%\Audian\cache</code> — the rewriting engine's evaluated instructions and examples (never your dictations); safe to delete

12 Troubleshooting

Symptom	What to do
Notification “Shortcut conflict”	Another app owns that shortcut. Pick another under Shortcut & recording (the editor checks availability live).
“No speech” after speaking	Check the microphone on Microphone & sound (Test microphone). Raise input gain or voice-detection sensitivity, or lower the silence threshold (the yellow mark on the meter).
Stops before I'm finished (hands-free)	Choose <i>Patient</i> or raise the pause tolerance; keep smart detection on.
Text copied instead of inserted	The target was closed or runs as administrator. Press Ctrl + V . To dictate into admin apps, run Audian as administrator too.
Nothing pastes in a specific app	General › Method › Simulate typing.
One slower dictation after changing mode, language or custom instructions	Normal: the rewriting engine prepares the new prompt once (~3.5 s) and caches it; later dictations restore it in milliseconds.
Back-to-back dictations feel slow on an older PC	Set “Keep the model in memory” to 1–5 minutes so the engines stay loaded between dictations.
Wrong language recognised	Speech recognition › Languages you speak, or fix the spoken language.
Antigravity errors	Make sure Antigravity is installed and signed in; the notification names the cause (sign-in, quota, network, model). Audian falls back to basic cleanup.
Names spelled wrong	Add them to Speech recognition › Custom vocabulary.

Diagnostics

```
audian.exe --transcribe recording.wav [--mode ai_prompt] [--provider antigravity] [--output en] > out.txt
audian.exe --rewrite-file lines.txt [--mode natural]      (prefix a line with "pt|" for Portuguese)
audian.exe --list-devices > devices.txt
audian.exe --reload                                     (re-read config.toml after manual edits)
audian.exe --export-art art-out                          (logo + recording-indicator images, every theme)
audian-stt.exe --bench ggml-small-q8_0.bin clip16k.wav 0 auto en,pt
audian-llm.exe --bench Qwen3.5-2B-Q4_K_M.gguf - "um so I want to fix the the bug"
```

Set the environment variable `AUDIAN_LOG=debug` for verbose logs.

13 Configuration reference

`%APPDATA%\Audian\config.toml` — everything is editable in the app; manual edits are applied with *Reload configuration* (tray > More) or `audian.exe --reload`. Missing keys use defaults.

Section · key	Default	Meaning
general · start_with_windows / notifications	false / true	Startup entry; error notifications
general · sounds / sound_volume	true / 0.5	Start/stop chimes
hotkey · shortcut / mode	"Ctrl+Shift+Space" / "hybrid"	hybrid · hold · toggle
hotkey · escape_cancels / paste_last_shortcut	true / "Ctrl+Shift+Alt+V"	Esc cancel; re-paste shortcut (" disables)
auto_stop · enabled / pause_secs	true / 2.0	Hands-free auto-stop and pause tolerance
auto_stop · smart / unfinished_extra_secs / no_speech_secs	true / 1.5 / 8.0	Sentence-aware waiting
audio · device_id / device_name	"" (system default)	Microphone
audio · gain / vad_sensitivity / silence_threshold / max_duration_secs	1.0 / 0.5 / 0.006 / 300	Input processing
transcription · language / auto_languages / vocabulary	"auto" / [] / []	Recognition language(s), glossary
transcription.whisper · model / threads / keep_loaded_minutes	"ggml-small-q8_0.bin" / 0 / 0	0 threads = auto; 0 minutes = unload right after use
processing · enabled / mode / provider	true / "natural" / "local_llm"	Rewriting on/off, default mode, provider (local_llm · antigravity · rules)
processing · custom_instructions / output_language / fallback_to_rules	"" / "same" / true	Extra instructions; translation target
processing.local_llm · model / threads / keep_loaded_minutes / max_output_tokens / timeout_secs	"Qwen3.5-2B-Q4_K_M.gguf" / 0 / 0 / 768 / 45	Local rewriting engine
processing.antigravity · executable / model / timeout_secs	"" / "gemini-3.8-flash-low" / 30	Cloud rewriting
profiles · enabled / rules[]	true / 5 rules	Per-app modes (name, enabled, mode, apps, title_keywords)
insertion · method / restore_clipboard / restore_delay_ms	"paste" / true / 600	How text is inserted
insertion · terminal_single_line / trailing_space	true / false	Formatting
overlay · position / show_label / style	"auto" / true / "dark"	auto · bottom_center · top_center; dark · light

appearance · theme	"ivory"	ivory · violet · ocean · mint · sunset · rose
privacy · history / history_limit / save_recordings / log_transcripts	true / 200 / false / false	Stored data

14 Building & publishing

Requirements: Rust (MSVC toolchain), Visual Studio 2022 Build Tools with *C++ build tools*, *Windows 11 SDK* and *C++ Clang Compiler for Windows*, plus CMake.

```
.\scripts\build.ps1      # app and engines → target\release
.\scripts\package.ps1   # + installer      → dist\Audian-Setup-1.0.0-x64.exe (prints size and SHA-256)
cargo test --release     # unit tests (config, IPC, VAD, auto-stop, hotkeys, DSP, rules, sanitizer)
```

Crate	Role
audian	Tray app + Audian window
audian-stt / audian-llm	Speech and rewriting engines
audian-setup	Installer/uninstaller; embeds the app binaries (deflate-compressed)
audian-common	Config, model catalog, downloader, history, IPC, logging, WAV
audian-ui / audian-art / audian-build	Design system · vector artwork · icon/manifest/version embedding

Before publishing

The installer is not code-signed, so Windows SmartScreen will warn users downloading it (“Windows protected your PC”). Signing it with a code-signing certificate (or distributing through the Microsoft Store / winget) removes the warning. Models are downloaded from Hugging Face on the user's machine — check the model licences (Whisper: MIT; Qwen 3.5: Apache-2.0; Llama 3.2: Llama Community License) if you redistribute them.